UNIVERSITÀ DEGLI STUDI DI MILANO

GPU COMPUTING PROJECT

Bitonic Sort using Numba on NVIDIA GPUs

Author:

MANUELE LUCCHI

08659A

ACADEMIC YEAR 2022/2023

1 Abstract

This documents describe the performances of an existing algorithm, the Bitonic Sort, implemented using Python and Numba, comparing it with a serial approach and CUDA C implementation.

Contents

2   Introduction

The Bitonic Mergesort [1] is an algorithm created in 1968 by Ken Batcher [2].
The elements to choose for the comparison are indipendent from the value (it's
not data-dependant) therefore is well suited for parallel processing.

Over the years there are been many implementations using low level languages
such as C and even high level such as Java or Python, but regarding the paral-
lel variant, only C/C++ examples can be found on the web (using OpenMP [3]
and CUDA C [4]) The CUDA community developed an abstraction layer over the
CUDA runtime to be able to execute CUDA kernels from python with little-to-no
overhead using Numba [5].

Numba is not a CUDA specific framework, it's a generic parallelization abstrac-
tion that allows to compile Python functions into machine code, while integrating
with numpy [6] to have access to a powerful tensor manipulation library; but of
course its primary target is CUDA.

The scope of the paper is to implement the algorithm using Numba and compare
it with different implementations

3   Algorithm

The algorithm is based on the concept of Bitonic Sequence [1], a sequence with
$x_0 \leq ... \leq x_k \geq x... \geq x_{n-1}$ for some k, $0 \leq k \leq n$, meaning that there are two
subsequences sorted in opposite directions.

By using a sorting network, we can create a Bitonic Sequence from any sequence
and then merging them to obtain the final sorted sequence, so the algorithm has
two phases.

This algorithm has an asymptotic complexity of $O(nlog(n)^2)$, the same of the
odd-even mergesort [7] and shellsort [8]

3.1   Step 1: Bitonic Sort

To create a Bitonic Sequence we need to build a sorting network. This step has log_2n phases (with the last one being the second step), each one with i stages, where i is the current phase starting from 1. By sorting each pair of elements in the sequence in different directions pairwise (using the so called "comparers"), we obtain a sequence full of bitonic subsequences. We can then at each phase double the size of these subsequences and half their number.

As we can see in picture one, at each stage of a phase the elements in the comparers get closer and in each phase a bitonic subsequence is sorted. This structure is known as butterfly network [9]

3.2   Step 2: Bitonic Merge

The last step is a variation of the first one, where we only have a single bitonic sequence and sort the two subsequences with the comparers oriented in the same direction, resulting in the sorted sequence.
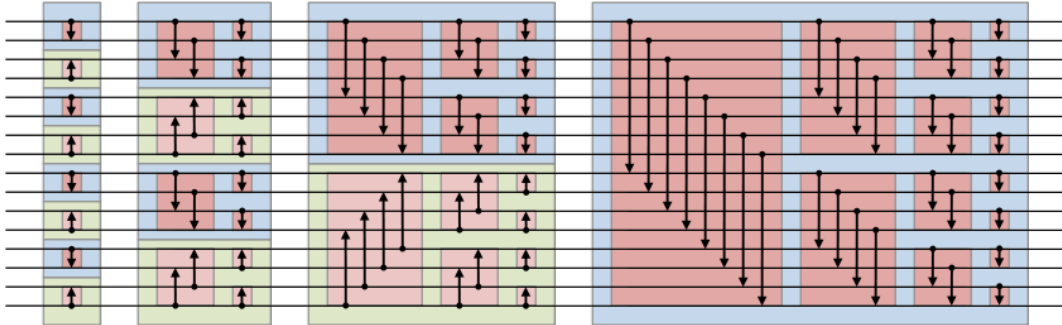

Figure 1: The structure of a bitonic sorting network

4   Implementation

The scope of the implementation is to use Numba to measure its efficiency compared to serial python algorithms and a standard CUDA C++ version.

The other versions are inspired by the course material and by the Wikipedia iterative implementation, with some changes to support a generic sorting direction

(and of course, translated to python).

The Python code consists of 3 functions

4.1   The host code

A function named *bitonic_sort* that accepts the initial data, its size and the direction. This function loops two indices, doubling the outer one each external iteration and halfing the inner one each internal iteration. This represents the $p = log_2 n$ phases and $s = log_2 p$ stages for each phase. Inside the two loops, the kernel function is called forwarding the data, the direction and passing the current phase and stage.

4.2   The kernel code

The *global* code, called by the host and executed in the device decides which pair of elements to compare is represented by the function *bitonic_kernel*.
It's decorated with the *@cuda.jit* decorator, that signales Numba that it's a function to compile to CUDA instructions.

4.3   The device code

4.4   Sequences of length not power of 2

5   Benchmark and Profiling

5.1   Environment

The benchmark and profiling environment is represented by a computer with an AMD Ryzen 7 5800H CPU (8 Cores, 16 Threads), 16 GB of RAM and an NVIDIA GeForce RTX 3050 Ti for Laptops. Basic specifications comprise **20 SMs** for a total of **2560 CUDA cores**, 20 RT cores and 80 Tensor cores. Clock speeds up to 1695 MHz boost at 80W, or as low as 1035 MHz for the 35W configuration,

giving us quite a wide power range of 35 to 80W. It has 4GB of GDDR6 memory on a 128-bit bus, clocked at 12 Gbps. The **Compute Capability** Version is 8.6.

5.2 JIT Compilation

The CUDA runtime executes functions, called kernels, compiled for NVIDIA GPUs instruction set using the NVCC Compiler [10] and this is a different approach compared to Python, which is a **runtime interpreted language**.

To solve this issue in the fastest way without changing paradigm was to compile **Just In Time** the python kernel into machine code, **caching it** and then executing it calling the CUDA Runtime Library. This of course comes with a delay in the first execution of a kernel (that needs to be compiled) and since the benchmarks would be cached except for the first one, the JIT compilation time has been considered separately. On the kernel function defined for the Bitonic Sort, the compilations on the environment takes around **0.3 seconds**, so it's noticeable if you want to compare with a sequential approach or a CUDA C version with a cold start.
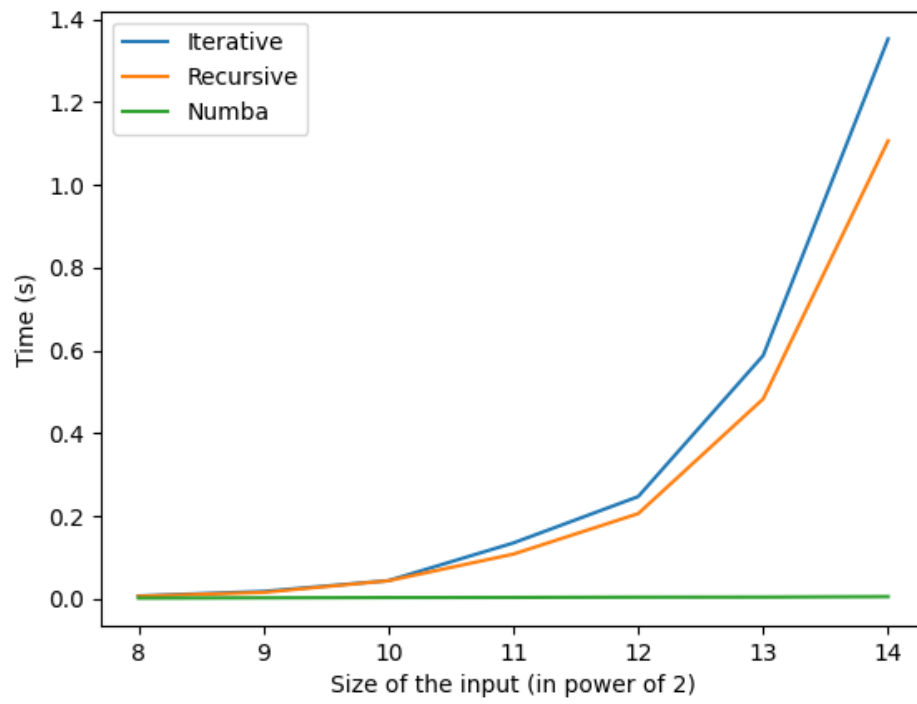
**Figure 2:** Time comparison between different implementations

Figure 3: Time comparison between parallel implementations on different plat-
forms

5.3 Serial and Parallel approach

5.4 CUDA C and Numba

Size (Power of 2)	CPU Recursive	CPU Iterative	GPU
8	10 ms	10 ms	10 ms
9	10 ms	10 ms	10 ms
10	10 ms	10 ms	10 ms
11	10 ms	10 ms	10 ms
12	10 ms	10 ms	10 ms
13	10 ms	10 ms	10 ms
14	10 ms	10 ms	10 ms
15	10 ms	10 ms	10 ms
16	10 ms	10 ms	10 ms
17	10 ms	10 ms	10 ms
18	10 ms	10 ms	10 ms
19	10 ms	10 ms	10 ms

Size (Power of 2)	Cuda C++	Python Numba
15	10 ms	10 ms
16	10 ms	10 ms
17	10 ms	10 ms
18	10 ms	10 ms
19	10 ms	10 ms
20	10 ms	10 ms
21	10 ms	10 ms
22	10 ms	10 ms
23	10 ms	10 ms
24	10 ms	10 ms
25	10 ms	10 ms
26	10 ms	10 ms
27	10 ms	10 ms

5.5   Profiling Results

Using the profiling tool *NCU* on a kernel execution, we obtained the following results:

5.5.1   GPU Speed Of Light Throughput

Metric Name	Metric Unit	Metric Value
DRAM Frequency	cycle/nsecond	5,75
SM Frequency	cycle/nsecond	1,17
Elapsed Cycles	cycle	2.757.034
Memory Throughput	%	92,97
DRAM Throughput	%	92,97
Duration	msecond	2,35
L1/TEX Cache Throughput	%	19,44
L2 Cache Throughput	%	29,83
SM Active Cycles	cycle	2.746.869,40
Compute (SM) Throughput	%	16,16

5.5.2   Launch Statistics

Metric Name	Metric Unit	Metric Value
Block Size		256
Function Cache Configuration		CachePreferNone
Grid Size		131.072
Registers Per Thread	register/thread	16
Shared Memory Configuration Size	Kbyte	8,19
Driver Shared Memory Per Block	Kbyte/block	1,02
Dynamic Shared Memory Per Block	byte/block	0
Static Shared Memory Per Block	byte/block	0
Threads	thread	33.554.432
Waves Per SM		1.092,27

5.5.3 Occupancy

Metric Name	Metric Unit	Metric Value
Block Limit SM	block	16
Block Limit Registers	block	16
Block Limit Shared Mem	block	8
Block Limit Warps	block	6
Theoretical Active Warps per SM	warp	48
Theoretical Occupancy	%	100
Achieved Occupancy	%	83,52
Achieved Active Warps Per SM	warp	40,09

6 Conclusion

References

1. Einstein, A. Zur Elektrodynamik bewegter Körper. (German) [On the electrodynamics of moving bodies]. *Annalen der Physik* **322,** 891–921 (1905).

2. Knuth, D. *Knuth: Computers and Typesetting* http://www-cs-faculty.stanford.edu/~uno/abcde.html. (accessed: 01.09.2016).

3. Knuth, D. *Knuth: Computers and Typesetting* http://www-cs-faculty.stanford.edu/~uno/abcde.html. (accessed: 01.09.2016).

4. Knuth, D. *Knuth: Computers and Typesetting* http://www-cs-faculty.stanford.edu/~uno/abcde.html. (accessed: 01.09.2016).

5. Knuth, D. *Knuth: Computers and Typesetting* http://www-cs-faculty.stanford.edu/~uno/abcde.html. (accessed: 01.09.2016).

6. Knuth, D. *Knuth: Computers and Typesetting* http://www-cs-faculty.stanford.edu/~uno/abcde.html. (accessed: 01.09.2016).

7. Knuth, D. *Knuth: Computers and Typesetting* http://www-cs-faculty.stanford.edu/~uno/abcde.html. (accessed: 01.09.2016).

8. Knuth, D. *Knuth: Computers and Typesetting* `http://www-cs-faculty.stanford.edu/~uno/abcde.html`. (accessed: 01.09.2016).

9. Knuth, D. *Knuth: Computers and Typesetting* `http://www-cs-faculty.stanford.edu/~uno/abcde.html`. (accessed: 01.09.2016).

10. Knuth, D. *Knuth: Computers and Typesetting* `http://www-cs-faculty.stanford.edu/~uno/abcde.html`. (accessed: 01.09.2016).